

CHAPTER-1

INTRODUCTION

1. INTRODUCTION

1.1 OVERVIEW:

Music genre prediction has become a prominent research area in the field of predictive analytics. The ability to automatically classify music into different genres has numerous applications, ranging from music recommendation systems to music analysis and understanding user preferences.

This project aims to develop an accurate music genre prediction model using various classification algorithms. By extracting meaningful features from audio data and applying advanced machine learning techniques, we can enhance the understanding and categorization of music based on its acoustic properties.

The goal of this project is to use the implementation and practical application of predictive analytics using multiple classification algorithms in order to identify the best classification or regression algorithm and use it to classify audio files using various concepts such as audio processing and other necessary functionalities. This will help create a model that can predict the genre of an audio file and this can be used in various scenarios such as personalized music suggestions, automated creation of personalized playlists in music streaming services, understanding listening patterns of individual users, and enhancing the overall music listening experience for all users.

The vision of this project is to establish the basis for understanding the best types of models that can be applied by analyzing accuracy measurements and choosing the most suitable algorithm. It also involves visualizing various aspects of the datasets as well as algorithm results. All of this provides valuable insight which can be used to make various inferences that prove to be fruitful in generating high quality analysis results.

1.2 STATEMENT OF THE PROBLEM

The problem addressed in this project is the accurate classification of music genres using predictive analytics. The challenge lies in effectively extracting relevant audio features that capture the distinct characteristics of different genres. Additionally, selecting the most suitable classification algorithm and evaluating its performance pose further challenges. The project seeks to overcome these obstacles and develop a robust model capable of accurately predicting music genres.

1.3 MOTIVATION

The motivation behind this project stems from the increasing demand for personalized music recommendation systems and the need for better music genre classification. By accurately predicting music genres, we can provide users with tailored music recommendations based on their preferences. Moreover, music genre prediction plays a vital role in music analysis, allowing us to gain insights into musical trends, patterns, and the evolution of genres over time.

1.4 CHALLENGES

Several challenges arise in the field of music genre prediction. One of the primary challenges is the inherent subjectivity and ambiguity associated with defining music genres. Genres can have overlapping characteristics and evolve over time, making it difficult to establish clear boundaries.

Another challenge is the variability in audio recordings, including differences in audio quality, instrumentation, and production techniques. Handling these variations and extracting robust features that capture genre-specific information is a significant challenge.

Additionally, the availability of labeled datasets with a diverse range of genres poses a challenge for training accurate models. Balancing the dataset to ensure sufficient representation of each genre is crucial for avoiding bias and achieving reliable predictions.

1.5 APPLICATIONS

The applications of music genre prediction using predictive analytics are diverse and span various domains. One key application is in the field of music recommendation systems, where accurate genre classification enables personalized music recommendations based on user preferences. It also finds utility in music streaming platforms, allowing for better categorization and organization of music libraries.

Music genre prediction can facilitate music analysis, such as studying the evolution of genres over time or identifying trends and patterns within specific genres. Additionally, it has potential applications in content creation, assisting artists and producers in exploring new genres and creating music with specific genre characteristics.

1.6 ORGANIZATION OF THE REPORT

This report is organized into five chapters, each addressing specific aspects of the project. Chapter 1 introduces the project, including an overview, problem statement, motivation, challenges, applications, and the organization of the report. Chapter 2 presents a literature survey, exploring previous work related to music genre prediction using predictive analytics. It includes a comprehensive review of relevant textbooks and their applications in this domain.

Chapter 3 focuses on the methodology employed in this project. It includes an overview of the proposed system architecture, a diagram illustrating the workflow, and detailed explanations of the steps involved. It also covers hardware and software requirements, programming languages, and processing libraries used in the project. Moreover, it provides an in-depth description of the GTZAN dataset, including its history, applications, and preprocessing techniques applied.

In Chapter 3, each algorithm used in the project is thoroughly described. This includes an overview of the algorithm, its general implementation for model fitting, and

specific details of its implementation within the music genre prediction context. The advantages and disadvantages of each algorithm are also discussed.

Chapter 4 presents the experiments and results obtained in the project. It provides a step-by-step breakdown of the project implementation, including data exploration, visualization techniques, dataset splitting, preprocessing, model training, and evaluation. Code snippets and annotations are provided to enhance understanding. The chapter also includes visual representations of the GTZAN dataset before and after processing, as well as a detailed explanation of each classification algorithm, along with accuracy scores and classification reports. The use of confusion matrices for result comparison is highlighted.

Chapter 5 concludes the report by summarizing the findings and contributions of the project. Future research directions and potential improvements are suggested. The report concludes with a comprehensive list of references used throughout the project, including relevant literature and online sources related to music genre prediction, GTZAN dataset, and machine learning algorithms.

CHAPTER-2

LITERATURE SURVEY

2. LITERATURE SURVEY

2.1 OVERVIEW:

The literature survey provides an overview of existing research and advancements in the field of music genre prediction using predictive analytics. It aims to establish a foundation of knowledge and identify key approaches and techniques employed in previous studies. By reviewing relevant textbooks, we gain insights into the theoretical and practical aspects of music genre classification.

2.2 LITERATURE REVIEW:

2.2.1 "Music Information Retrieval: Recent Developments and Applications"

AUTHORS: Zhiyao Duan, Bochen Li, and Changsheng Xu

This source discusses the application of music information retrieval techniques in various domains, including music genre classification. It provides an overview of feature extraction methods, such as spectrogram analysis, mel-frequency cepstral coefficients (MFCC), and rhythmic features. The book also explores different machine learning algorithms commonly used in music genre prediction, such as decision trees, support vector machines (SVM), and neural networks. It presents case studies and practical examples to demonstrate the effectiveness of these techniques in genre classification tasks.

2.2.2 "Pattern Recognition and Machine Learning"

AUTHOR: Christopher M. Bishop

Here, we delve into the fundamentals of pattern recognition and machine learning, which are essential components of music genre prediction. It covers topics such as Bayesian decision theory, classification algorithms (including k-nearest neighbors, decision trees, and support vector machines), and model evaluation metrics. The book provides a solid theoretical foundation for understanding the algorithms used in music genre classification and their mathematical principles. It also offers insights into

feature selection techniques, dimensionality reduction methods, and ensemble learning approaches.

2.2.3 "Introduction to Machine Learning"

AUTHOR: Ethem Alpaydin

This book offers an introduction to machine learning techniques and their applications. It covers the basics of supervised and unsupervised learning algorithms, including decision trees, support vector machines, k-nearest neighbors, and Gaussian Naive Bayes. The book provides a comprehensive overview of the algorithms' underlying concepts, their strengths, and limitations. It also discusses feature engineering, model evaluation, and the importance of dataset preparation for accurate predictions. The textbook includes practical examples and case studies that illustrate the application of machine learning in various domains, including music genre classification.

2.2.4 "Music Data Analysis: Foundations and Applications"

AUTHOR: Tao Li, Mitsunori Ogihara, and George Tzanetakis

Focuses specifically on music data analysis, including music genre classification. It covers key techniques and methodologies used in analyzing music data, such as feature extraction, data preprocessing, and classification algorithms. The book explores various feature representations, including audio-based features, symbolic representations, and metadata-based features. It discusses the challenges and considerations specific to music data analysis, such as the extraction of meaningful features from audio signals and the interpretation of genre boundaries. The textbook includes case studies and practical examples to demonstrate the application of music genre classification techniques in real-world scenarios.

2.2.5 "Music Information Retrieval: A Practical Guide for Audio Signal Processing" AUTHOR: Meinard Müller

Serves as a practical guide to music information retrieval, covering various aspects of audio signal processing for music analysis tasks. It introduces fundamental concepts, such as spectral analysis, pitch estimation, and rhythm analysis, which are

relevant to music genre classification. The book explores feature extraction techniques specific to music signals, including timbral features, harmonic features, and rhythmic features. It also discusses machine learning algorithms commonly used in music genre classification, such as decision trees, random forests, and support vector machines. Practical examples and code snippets are provided to aid understanding and implementation.

2.2.6 "Music Genre Classification: A Comparison of Popular Approaches"

AUTHOR: Giorgi Chaladze

This paper compares different approaches for music genre classification, including feature extraction techniques, classification algorithms, and evaluation metrics. It provides insights into the performance of various methods and their applicability in real-world scenarios.

2.2.7 "Music Genre Classification: A Multiclass Problem"

AUTHOR: Giovanni Scagliola, et al.

This conference paper focuses on the multiclass classification problem in music genre classification. It explores different feature representations, classification algorithms, and evaluation methodologies. The study provides a comprehensive overview of the challenges and techniques in the field.

2.2.8 "Music Genre Classification: A Case Study with Top-100 Songs in Ten Genres"

AUTHOR: Bob L. Sturm, et al.

This research article presents a case study on music genre classification using a dataset of top-100 songs from ten different genres. It discusses the challenges associated with genre classification, feature extraction methods, and classifier selection. The study also includes an evaluation of the classification performance using various algorithms.

These literature sources and documents, among others, serve as valuable references for understanding the theoretical foundations, practical methodologies, and applications of music genre prediction using predictive analytics. They provide insights into feature extraction techniques, classification algorithms, and evaluation metrics used in the field. By leveraging the knowledge and findings from previous studies, this project aims to build upon existing research and contribute to the advancement of music genre prediction algorithms.

CHAPTER-3

METHODOLOGY

3. METHODOLOGY:

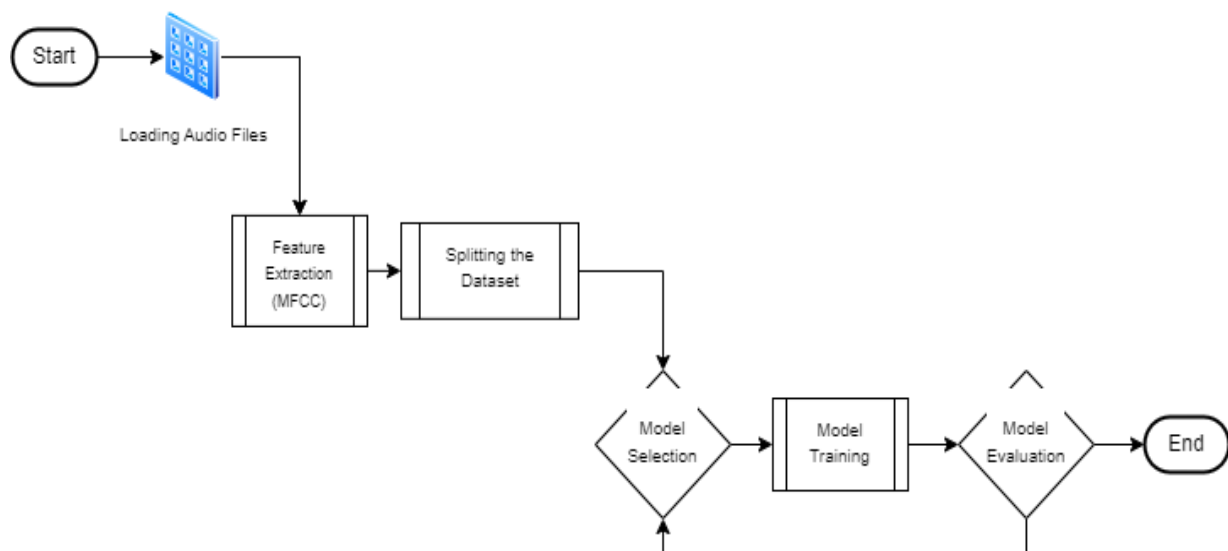
3.1 OVERVIEW:

This chapter outlines the methodology employed in the project for music genre prediction using predictive analytics. It provides an overview of the proposed system architecture, highlighting the key components and their interactions. The chapter also discusses the hardware and software requirements, programming languages, and processing libraries utilized in the project.

We aim to produce a model that runs through multiple audio files, trains a model, and produces results by generating a synthetic data set of features that are extracted from the processing of the audio files, and running them through algorithms to predict the genre.

3.2 ARCHITECTURE OF PROPOSED SYSTEM / SYSTEM DESIGN:

The proposed system for music genre prediction comprises several interconnected modules that work together to achieve accurate genre classification. The system architecture is depicted in Figure 1.



The main modules of the system design include:

3.2.1 Data Collection and Preprocessing

This module focuses on gathering a diverse dataset of music audio files from various genres. The collected data undergoes preprocessing steps, such as audio normalization, noise removal, and format conversion, to ensure consistency and quality.

3.2.2 Feature Extraction

The feature extraction module transforms the raw audio data into a representative feature space that captures genre-specific information. Various techniques are applied to extract features, such as spectral features, temporal features, and tonal features. These features serve as inputs to the classification algorithms.

3.2.3 Model Training

In this module, classification models are trained using the extracted features and corresponding genre labels. Different machine learning algorithms, including decision trees, support vector machines, and neural networks, are utilized for model training.

3.2.4 Evaluation

The trained models are evaluated using appropriate evaluation metrics, such as accuracy, precision, recall, and F1 score. The metrics of each model are then compared, which helps in identifying and choosing the best algorithm that produces the most accurate classification results.

3.2.5 Genre Prediction

Once the models are trained and evaluated, the genre prediction module takes in new, unseen audio samples and applies the trained models to predict their respective genres. The predicted genres are outputted as the results of the system.

3.3 HARDWARE AND SOFTWARE REQUIREMENTS:

3.3.1 Hardware Requirements

- (i). A computer system with sufficient processing power and memory to handle the dataset size and computational requirements of the machine learning algorithms.
- (ii). Adequate storage capacity to store the dataset, feature vectors, and trained models.

3.3.2 Software Requirements

- (i) Operating System: Any modern operating system, such as Windows, macOS, or Linux.
- (ii) Python: A programming language used for data processing, feature extraction, model training, and prediction.
- (iii) Integrated Development Environment (IDE): A suitable Python IDE, such as Anaconda, Jupyter Notebook, or PyCharm, for code development and execution.
- iv) Processing Libraries: Relevant Python libraries, including NumPy, Pandas, SciPy, scikit-learn, and librosa, for data manipulation, feature extraction, machine learning, and audio signal processing tasks.

3.4 LANGUAGE USED

Python is a powerful and versatile programming language that was created by Guido van Rossum between 1985 and 1990. It is a high-level, interpreted, interactive, and object-oriented language that aims to be highly readable. Unlike other languages, Python uses English keywords instead of punctuation and has a simpler syntax. It is widely used for various applications, including implementing bank payment simulations for fraud detection. Python's design philosophy emphasizes code reliability, and its concise syntax allows programmers to express concepts in fewer lines of code compared to languages like C++ or Java. Python supports multiple programming paradigms, including

object-oriented, imperative, and functional programming styles. It also features a dynamic type system, automatic memory management, and a comprehensive standard library.

Python's popularity is due in part to its ease of use, making it an ideal language for beginners. It provides a gentle learning curve for novice programmers, enabling them to develop a wide range of applications, from simple text processing to web browsers and games. Python's interactive nature allows programmers to directly interact with the interpreter, making it convenient for testing and debugging code. Additionally, Python's object-oriented capabilities enable the encapsulation of code within objects, promoting modularity and reusability.

In practical terms, when working on a project that utilizes Python, it is common to use packages. A package is a hierarchical file directory structure that defines a single Python application environment, consisting of modules and sub-packages. These packages can be installed to extend Python's functionality and access additional libraries or frameworks. The availability of Python interpreters for various operating systems further enhances its versatility and makes it accessible to a wide range of developers.

Overall, Python's combination of readability, versatility, and user-friendliness has contributed to its widespread adoption as a programming language for diverse applications, from small-scale scripts to large-scale projects.

Python's rich library ecosystem, including NumPy, Pandas, and Scikit-learn, makes it a popular choice for predictive analytics, as it also provides highly accurate predictions and generates efficient classification models.

3.5 GTZAN DATASET

The GTZAN dataset is a widely used benchmark dataset in the field of music genre classification. It contains 1,000 audio files, each approximately 30 seconds long, representing ten different music genres. The genres include blues, classical, country, disco, hip-hop, jazz, metal, pop, reggae, and rock. The dataset provides a balanced distribution of samples across genres, making it suitable for training and evaluating music genre prediction models.

The GTZAN dataset has been extensively utilized in research to assess the performance of various classification algorithms. Its popularity stems from its wide availability, established ground truth labels, and diverse representation of music genres. Preprocessing techniques, such as audio format conversion, audio normalization, and feature extraction, are often applied to the dataset before using it for model training and evaluation.

3.6 PACKAGES USED

The project employs several packages to facilitate the implementation of various tasks. The following packages are utilized:

1. NumPy: A fundamental package for scientific computing in Python, providing support for efficient numerical operations and array manipulation.

2. Pandas: A powerful library for data manipulation and analysis. It enables efficient handling of structured data, such as loading and preprocessing the dataset, and performing exploratory data analysis.

3. SciPy: A library that provides scientific and mathematical functionality in Python. It includes modules for signal processing, statistical operations, and optimization, which are useful for tasks such as audio signal processing and feature extraction.

4. scikit-learn: A comprehensive machine learning library in Python. It provides a wide range of classification, regression, and evaluation algorithms, including decision trees, support vector machines, and evaluation metrics for model assessment.

5. librosa: A Python library specifically designed for audio and music signal analysis. It offers various functionalities for audio feature extraction, such as spectrogram computation, mel-frequency cepstral coefficients (MFCC), and tempo estimation.

6. Matplotlib: A popular plotting library in Python for creating visualizations and graphs. It is utilized to generate data visualizations, such as spectrograms, waveform plots, and accuracy comparison charts.

7. Seaborn: A data visualization library built on top of Matplotlib. It provides a high-level interface for creating aesthetically pleasing statistical graphics and enhancing the visual presentation of data.

8. Jupyter Notebook: An interactive development environment that allows for the creation and sharing of documents containing live code, equations, visualizations, and explanatory text. It is used for code development, experimentation, and generating reports.

9. Anaconda: A distribution of Python and associated packages that simplifies package management and provides a comprehensive environment for scientific computing and data analysis.

3.7 DATA CLASSIFICATION

In the context of music genre prediction, several key terminologies are relevant. These include:

1. Classifier: A classifier is a machine learning algorithm or model that learns patterns from input data and assigns them to predefined categories or classes. It forms the core component of the music genre prediction system.

2. Classification Model: A classification model is the result of training a classifier on labeled data. It captures the learned patterns and relationships between input features and their corresponding output labels, enabling the prediction of the genre for new, unseen music samples.

3. Feature: A feature refers to a measurable characteristic or attribute extracted from the input data. In music genre prediction, features represent different aspects of the audio signal, such as spectral content, rhythmic patterns, or tonal characteristics. These features serve as input variables to the classification model.

4. Accuracy Measures: Accuracy measures are metrics used to assess the performance of a classification model. They provide insights into the model's ability to correctly predict the genre of music samples. Common accuracy measures include accuracy, precision, recall, and F1 score, which quantify different aspects of the classification performance.

5. Confusion Matrix: A confusion matrix is a table that is used to define the performance of a classification algorithm. A confusion matrix visualizes and summarizes the performance of a classification algorithm.

6. Multi-label Classification: classification task where each sample is mapped to a set of target labels (more than one class).

7. Regression Model: A regression model is a statistical model that estimates the relationship between one dependent variable and one or more independent variables using a line (or a plane in the case of two or more independent variables).

3.8 ALGORITHMS USED TO FIT THE MODEL

In this project, several classification and regression algorithms are utilized to fit the music genre prediction model. Each algorithm has its own advantages and disadvantages. The algorithms used in the project, along with their general descriptions and specific implementations, are as follows:

3.8.1 Decision Trees:

- **General Description:** Decision trees are hierarchical structures that recursively partition the input space based on the values of input features. They make decisions by traversing the tree from the root to the leaf nodes, where each leaf represents a class label.
- **Implementation:** The scikit-learn library provides a `DecisionTreeClassifier` class that implements decision tree-based classification. It allows for the specification of various parameters, such as the maximum depth of the tree and the criterion for splitting nodes.
- **Advantages:** Decision trees are easy to understand and interpret. The resulting tree structure can be visualized and easily explained to non-technical stakeholders.
- **Disadvantages:** Decision trees can be prone to overfitting, especially if the tree grows too deep or is not properly pruned. This can result in poor generalization to unseen data.

3.8.2 Random Forest Classifier:

- **General Description:** Random forests are an ensemble learning method that combines multiple decision trees. Each tree is trained on a different subset of the dataset, and the final prediction is obtained through majority voting or averaging the predictions of individual trees.
- **Implementation:** The `RandomForestClassifier` class from scikit-learn is used to implement random forest-based classification. It allows for controlling parameters like the number of trees in the forest and the maximum depth of each tree.

- **Advantages:** Random forests are an ensemble of decision trees, which leads to improved performance and robustness compared to a single decision tree. They can handle many features and provide estimates of feature importance.
- **Disadvantages:** Random forests can be computationally expensive and may require more memory compared to individual decision trees. They are also harder to interpret due to the ensemble nature.

3.8.3 Support Vector Machines:

- **General Description:** Support Vector Machines are powerful supervised learning algorithms that separate data points into different classes using hyperplanes in high-dimensional feature spaces. They aim to find the optimal hyperplane that maximizes the margin between different classes.
- **Implementation:** The scikit-learn library provides the SVC class for implementing SVM-based classification. It offers flexibility in choosing different kernels (linear, polynomial, radial basis function) and tuning hyperparameters for better model performance.
- **Advantages:** SVMs can effectively handle high-dimensional data and are particularly useful when the number of features is larger than the number of instances. They can also handle non-linear relationships through the use of kernel functions.
- **Disadvantages:** SVMs can be sensitive to the choice of kernel and its parameters. Tuning the kernel and regularization parameters can be time-consuming and require expertise. SVMs can also be computationally expensive, especially for large datasets.

3.8.4 K-Nearest Neighbors:

- **General Description:** K-Nearest Neighbors is a simple and intuitive algorithm that classifies data points based on the class labels of their neighboring data points. It measures the distance between instances in the feature space and assigns the class label based on the majority class among the K nearest neighbors.

- **Implementation:** The KNeighborsClassifier class in scikit-learn allows for implementing KNN-based classification. It enables customization of the number of neighbors (K) and the distance metric used for classification.
- **Advantages:** KNN is a non-parametric algorithm that does not assume any specific distribution of the data. It can handle complex decision boundaries and is easy to understand and implement.
- **Disadvantages:** KNN can be sensitive to the choice of the number of neighbors (K). Choosing an inappropriate K value may lead to overfitting or underfitting. It can also be computationally expensive, especially for large datasets.

3.8.5 Logistic Regression:

- **General Description:** Logistic regression is a statistical algorithm used for binary classification problems. It models the relationship between a dependent variable (binary outcome) and one or more independent variables by estimating the probabilities using a logistic function. It aims to find the best-fitting line that separates the two classes in the feature space.
- **Implementation:** The LogisticRegression class in scikit-learn provides an implementation of logistic regression for classification tasks. It supports various regularization techniques and allows customization of parameters like regularization strength and solver method.
- **Advantages:** Logistic regression provides interpretable coefficients that can indicate the direction and magnitude of the influence of each feature on the outcome. It can also handle both categorical and continuous features.
- **Disadvantages:** Logistic regression assumes a linear relationship between the features and the log-odds of the outcome, which may not hold in some cases. It may not perform well when the relationship is non-linear or when there are interactions between features.

3.8.6 Naïve-Bayes Classifier:

- **General Description:** The Naive Bayes classifier is a probabilistic algorithm based on Bayes' theorem with the assumption of feature independence. It predicts the class label of an instance by calculating the posterior probability of each class given the input features and selecting the class with the highest probability.
- **Implementation:** The GaussianNB, MultinomialNB, and BernoulliNB classes in scikit-learn provide implementations of different variants of the Naive Bayes classifier. GaussianNB assumes features follow a Gaussian distribution, MultinomialNB is suitable for discrete features, and BernoulliNB is designed for binary features.
- **Advantages:** Naive Bayes classifiers are computationally efficient and can handle high-dimensional data well. They perform particularly well when the independence assumption holds, and they require less training data compared to some other algorithms.
- **Disadvantages:** The assumption of feature independence may not hold true in real-world datasets, which can lead to suboptimal performance. Naive Bayes classifiers may struggle to capture complex relationships between features.

CHAPTER-4

EXPERIMENTS AND RESULTS

4. EXPERIMENTS AND RESULTS

We conducted a series of experiments to evaluate the performance of various classification models in the music genre prediction using predictive analytics. The experiments were carried out on the GTZAN dataset, which consists of audio samples from different music genres.

To preprocess the audio data, we employed the Mel-frequency cepstral coefficients (MFCC) technique. MFCC is a representation of the short-term power spectrum of an audio signal, which captures the spectral characteristics of the music. It provides a compact representation of the audio data by extracting relevant features such as timbre and texture.

The experimental process involved the following steps:

1. **Data Loading and Preprocessing:** We loaded the audio samples from the GTZAN dataset and applied MFCC feature extraction to convert the raw audio into a set of numerical features.
2. **Feature Extraction:** The MFCC algorithm computed the mel spectrogram, which represents the energy distribution across different frequency bands. It then applied a logarithmic transformation and performed a discrete cosine transform to obtain the MFCC coefficients.
3. **Data Splitting:** We divided the dataset into training and testing sets to assess the performance of the classification models. The training set was used to train the models, while the testing set was used for evaluation.
4. **Model Training and Evaluation:** We employed various classification algorithms, including K-Nearest Neighbors, Logistic Regression, Random Forest, Decision Tree, and Support Vector Machine (SVM), to train the models using the training data. Each model was evaluated using performance metrics such as accuracy, precision, recall, and F1 score.

5. Result Analysis: The classification results were analyzed and compared using a confusion matrix, which provided insights into the accuracy and misclassification of the models across different music genres.

By conducting these experiments, we gained valuable insights into the performance of different classification models in music genre prediction. The results demonstrated the effectiveness of the models in accurately classifying music samples into their respective genres. The experiments highlighted the potential of predictive analytics techniques in automating music genre identification, which can have applications in recommendation systems, music streaming platforms, and content categorization.

Overall, the experimental process allowed us to explore and evaluate the capabilities of various classification models in the context of music genre prediction, providing a foundation for further research and development in this field.

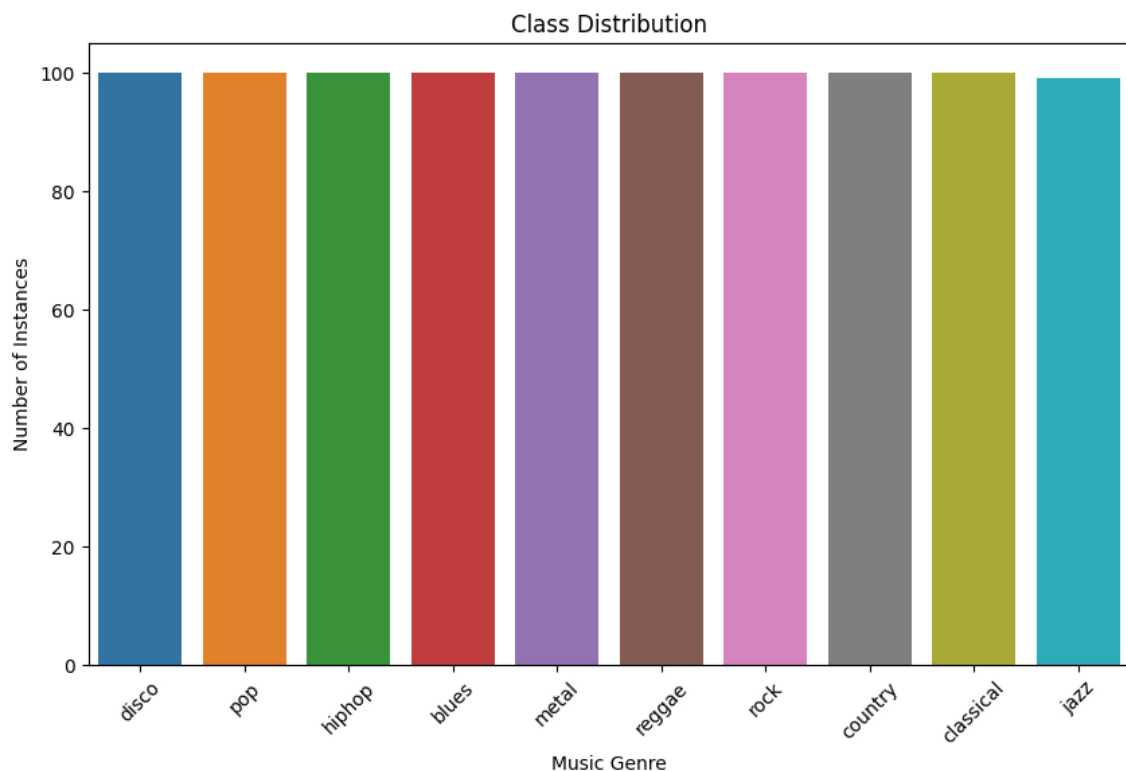
4.1 DATA EXPLORATION AND VISUALIZATION

Before delving into the experiment details, it is essential to perform data exploration and visualization to gain insights into the GTZAN dataset and understand its characteristics. This section presents additional code snippets for data exploration and visualization, as well as a description of the results.

Step 4.1.1: Class Distribution Visualization

One common aspect of data exploration is to examine the class distribution in the dataset. This can be achieved by visualizing the number of instances for each music genre. The following code snippet accomplishes this:

```
[5] #Class Distribution Visualization
genre_counts = df['label'].value_counts()
plt.figure(figsize=(10, 6))
sns.barplot(x=genre_counts.index, y=genre_counts.values)
plt.xlabel('Music Genre')
plt.ylabel('Number of Instances')
plt.title('Class Distribution')
plt.xticks(rotation=45)
plt.show()
```




Here, each genre has an equal number of class instances (audio files) which is 100.

Step 4.1.2: Feature Visualization

In addition to class distribution, it is beneficial to visualize the audio features across different genres. This can help identify potential patterns or differences that may exist between genres. One way to accomplish this is by creating box plots for each feature based on the genre. Here is an example of how this can be done:

```
[16] #Feature Visualization
plt.figure(figsize=(12, 8))
sns.boxplot(x='label', y=f'mfcc_{1}', data=df)
plt.xlabel('Music Genre')
plt.ylabel(f'MFCC {1}')
plt.tight_layout()
plt.show()
```



(MFCCs) extracted from the audio files. Ideally, each audio file has its own unique MFCCs and its visualization helps in identifying any variations or similarities between genres in terms of the extracted audio features.

4.2 EXPERIMENTAL PROCEDURE

The experiment involves the following steps:

Step 4.2.1: Splitting the Dataset

The GTZAN dataset is divided into training and testing sets using the `train_test_split` function from `scikit-learn`. The splitting is performed with a test size of 0.2 and a random state of 42 to ensure reproducibility.

```
# Split the dataset into training and testing sets, while maintaining genre balance
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

This code snippet utilizes the `train_test_split` function from the `scikit-learn` library to split the features (X) and labels (y) into training and testing sets. The `test_size` parameter specifies the proportion of the dataset to be used for testing (20% in this

case), and the `random_state` parameter ensures that the data split remains consistent across different runs.

Step 4.2.2: Preprocessing the Data

Prior to training the models, it is necessary to preprocess the data. In the given code, the `extract_features` function is used to extract the MFCC features from the audio files. The extracted features are then stored in a data frame with corresponding genre labels.

```
[4] def extract_features(file_path):
    try:
        audio_data, sample_rate = librosa.load(file_path, res_type='kaiser_fast')
        mfccs = librosa.feature.mfcc(y=audio_data, sr=sample_rate, n_mfcc=40)
        mfccs_scaled = np.mean(mfccs.T, axis=0)
        return mfccs_scaled
    except Exception as e:
        print(f"Error encountered while processing {file_path}: {str(e)}")
        return None

data = []
labels = []

# Iterate through the main folder and its subfolders
for root, dirs, files in os.walk('/content/drive/MyDrive/gtzan_music'):
    for file_name in files:
        if file_name.endswith('.wav'):
            file_path = os.path.join(root, file_name)
            features = extract_features(file_path)
            if features is not None:
                data.append(features)
                labels.append(file_name.split('.')[0]) # Assuming the genre is in the filename
            else:
                print(f"Skipping {file_path} due to errors during feature extraction.")
```

This code snippet demonstrates the process of feature extraction from the audio files using the `extract_features` function which iterates and extracts labels from them.

GTZAN DATASET:

	mfcc_2	mfcc_3	mfcc_4	mfcc_5	mfcc_6	mfcc_7	mfcc_8	mfcc_9	mfcc_10	label
0	51.860497	-7.567608	21.507437	8.141378	18.118084	6.894575	6.569047	0.711365	11.159763	disco
1	64.491341	5.644417	2.701930	0.786340	6.165176	1.884115	1.974974	-1.810813	1.865811	disco
2	57.173725	21.160828	11.305264	18.118341	1.662812	0.345103	2.323979	-5.105269	2.747911	disco
3	52.320511	3.242855	43.939304	12.584174	24.506044	4.923503	15.306510	2.041092	10.386366	disco
4	94.025856	12.438059	16.629438	8.126040	15.663839	-1.410333	7.014441	0.959783	-1.499451	disco
...
994	100.832466	-1.431354	41.276501	-0.943615	23.378519	-11.509500	20.015314	-15.975561	13.942784	jazz
995	100.935074	17.103338	12.981271	16.059950	4.768146	4.571938	-3.135014	-4.152950	-2.297984	jazz
996	118.277527	31.093847	40.150215	13.094316	13.939775	-3.718762	4.747173	-9.887995	-3.066447	jazz
997	112.720093	9.759592	38.984005	3.941792	20.022207	-7.214859	15.278688	-13.678443	10.154262	jazz
998	102.871300	-1.664055	45.214291	-6.754359	25.362495	-14.371586	24.249865	-10.413199	11.150229	jazz


Step 4.2.3: Model Training and Evaluation

1. K-Neighbors Classifier

Accuracy: 0.495

Classification Report:

	precision	recall	f1-score	support
blues	0.60	0.60	0.60	20
classical	0.68	0.75	0.71	20
country	0.38	0.60	0.46	20
disco	0.43	0.45	0.44	20
hiphop	0.27	0.20	0.23	20
jazz	0.41	0.35	0.38	20
metal	0.71	0.60	0.65	20
pop	0.68	0.75	0.71	20
reggae	0.33	0.30	0.32	20
rock	0.44	0.35	0.39	20
accuracy			0.49	200
macro avg	0.49	0.49	0.49	200
weighted avg	0.49	0.49	0.49	200



EXPLANATION: Step 1: Import the KNeighborsClassifier class from the sklearn.neighbors library, which allows us to use the K-Nearest Neighbors algorithm.

Step 2: Initialize the KNeighborsClassifier with the n_neighbors parameter, specifying the number of neighbors to consider for classification.

Step 3: Fit the model to the GTZAN dataset using the fit method, which trains the classifier on the input data and their corresponding labels.

GTZAN Music Genre Classification

Step 4: Evaluating the algorithm, confusion matrix, precision, recall and f1 score are the most used metrics. The confusion_matrix and classification_report methods of the sklearn.metrics can be used to calculate these metrics.


Step 5: The result shows that our K-Neighbors classifier algorithm was able to classify the test set with 0.495 accuracy in the GTZAN dataset.

2. Logistic Regression

Accuracy: 0.565

Classification Report:

	precision	recall	f1-score	support
blues	0.42	0.55	0.48	20
classical	0.90	0.90	0.90	20
country	0.40	0.50	0.44	20
disco	0.41	0.45	0.43	20
hiphop	0.45	0.50	0.48	20
jazz	0.67	0.50	0.57	20
metal	0.76	0.80	0.78	20
pop	0.78	0.70	0.74	20
reggae	0.45	0.45	0.45	20
rock	0.55	0.30	0.39	20
accuracy			0.56	200
macro avg	0.58	0.57	0.57	200
weighted avg	0.58	0.56	0.57	200



EXPLANATION: Step 1: Import the LogisticRegression class from sklearn.linear_model library.

Step 2: Initialize the LogisticRegression model.

Step 3: Fit the model to the GTZAN dataset.

Step 4: Evaluate the algorithm using metrics such as confusion matrix, precision, recall, and F1 score.


Step 5: The Logistic Regression model achieved an accuracy of 0.565 on the GTZAN dataset.

3. Random Forest Classification

Accuracy: 0.6

Classification Report:

	precision	recall	f1-score	support
blues	0.58	0.55	0.56	20
classical	0.77	0.85	0.81	20
country	0.42	0.50	0.45	20
disco	0.41	0.45	0.43	20
hiphop	0.53	0.45	0.49	20
jazz	0.72	0.65	0.68	20
metal	0.69	0.90	0.78	20
pop	0.76	0.80	0.78	20
reggae	0.59	0.50	0.54	20
rock	0.50	0.35	0.41	20
accuracy			0.60	200
macro avg	0.60	0.60	0.59	200
weighted avg	0.60	0.60	0.59	200



EXPLANATION: Step 1: Import the RandomForestClassifier class from sklearn.ensemble library.

Step 2: Initialize the RandomForestClassifier model.

Step 3: Fit the model to the GTZAN dataset.

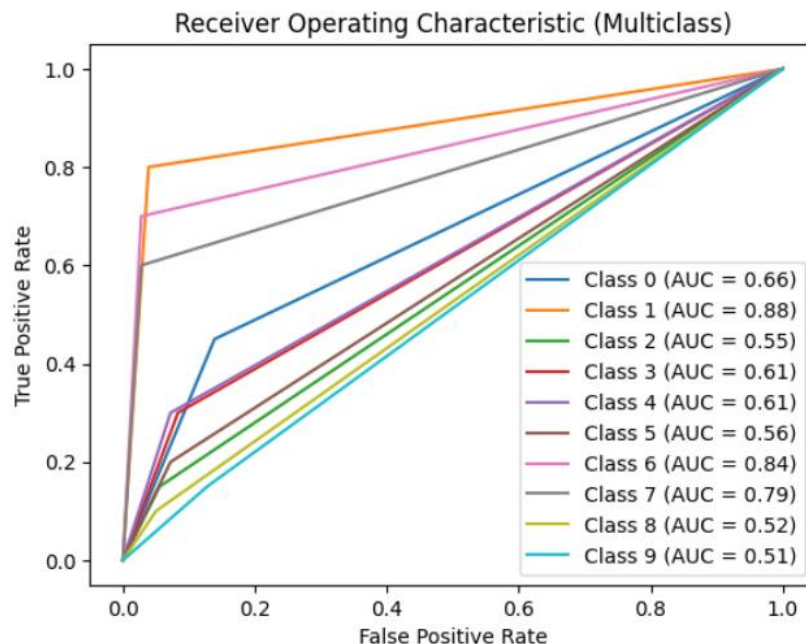
Step 4: Evaluate the algorithm using metrics such as confusion matrix, precision, recall, and F1 score.

Step 5: The Random Forest Classifier achieved an accuracy of 0.6 on the GTZAN dataset.

4. Decision Tree Classifier

Accuracy: 0.375
Classification Report:

	precision	recall	f1-score	support
blues	0.26	0.45	0.33	20
classical	0.70	0.80	0.74	20
country	0.23	0.15	0.18	20
disco	0.29	0.30	0.29	20
hiphop	0.32	0.30	0.31	20
jazz	0.24	0.20	0.22	20
metal	0.74	0.70	0.72	20
pop	0.71	0.60	0.65	20
reggae	0.18	0.10	0.13	20
rock	0.12	0.15	0.13	20
accuracy			0.38	200
macro avg	0.38	0.38	0.37	200
weighted avg	0.38	0.38	0.37	200



EXPLANATION: Step 1: Import the DecisionTreeClassifier class from sklearn.tree library.

Step 2: Initialize the DecisionTreeClassifier model.

Step 3: Fit the model to the GTZAN dataset.


Step 4: Evaluate the algorithm using metrics such as confusion matrix, precision, recall, and F1 score.

Step 5: The Decision Tree Classifier achieved an accuracy of 0.375 on the GTZAN dataset.

5. Support Vector Machine (SVM)

Accuracy: 0.435
Classification Report:

	precision	recall	f1-score	support
blues	0.25	0.10	0.14	20
classical	0.80	0.80	0.80	20
country	0.23	0.15	0.18	20
disco	0.25	0.15	0.19	20
hiphop	0.53	0.40	0.46	20
jazz	0.25	0.30	0.27	20
metal	0.55	0.90	0.68	20
pop	0.51	0.90	0.65	20
reggae	0.43	0.45	0.44	20
rock	0.21	0.20	0.21	20
accuracy			0.43	200
macro avg	0.40	0.44	0.40	200
weighted avg	0.40	0.43	0.40	200



EXPLANATION: Step 1: Import the SVC (Support Vector Classifier) class from the sklearn.svm library.

Step 2: Initialize the SVC model with the desired parameters.

Step 3: Fit the model to the GTZAN dataset.

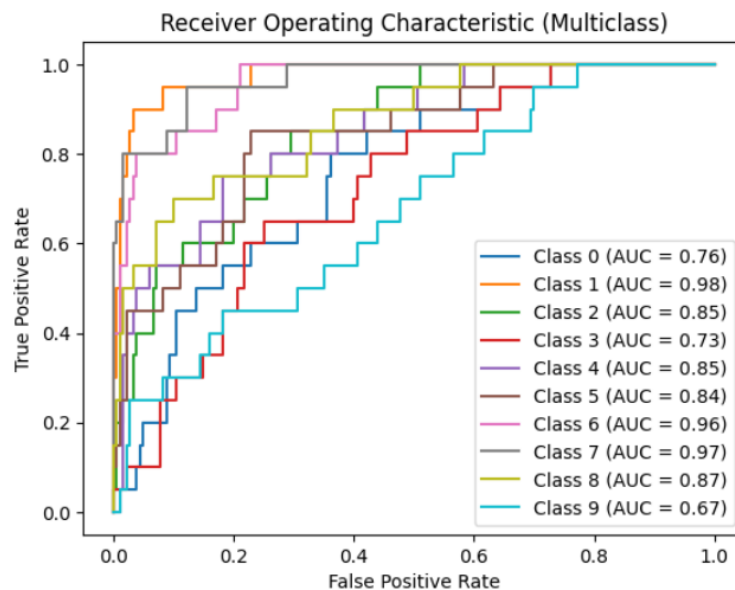
Step 4: Evaluate the algorithm using metrics such as confusion matrix, precision, recall, and F1 score.

Step 5: The Support Vector Classifier achieved an accuracy of 0.435 on the GTZAN dataset.

6. Naïve Bayes Classifier

Accuracy: 0.495
Classification Report:

	precision	recall	f1-score	support
blues	0.25	0.15	0.19	20
classical	0.77	0.85	0.81	20
country	0.42	0.50	0.45	20
disco	0.33	0.05	0.09	20
hiphop	0.69	0.45	0.55	20
jazz	0.56	0.45	0.50	20
metal	0.44	0.85	0.58	20
pop	0.44	0.85	0.58	20
reggae	0.65	0.55	0.59	20
rock	0.33	0.25	0.29	20
accuracy			0.49	200
macro avg	0.49	0.49	0.46	200
weighted avg	0.49	0.49	0.46	200



Step 1: Import the NaiveBayes classifier class from the sklearn.naive_bayes library.

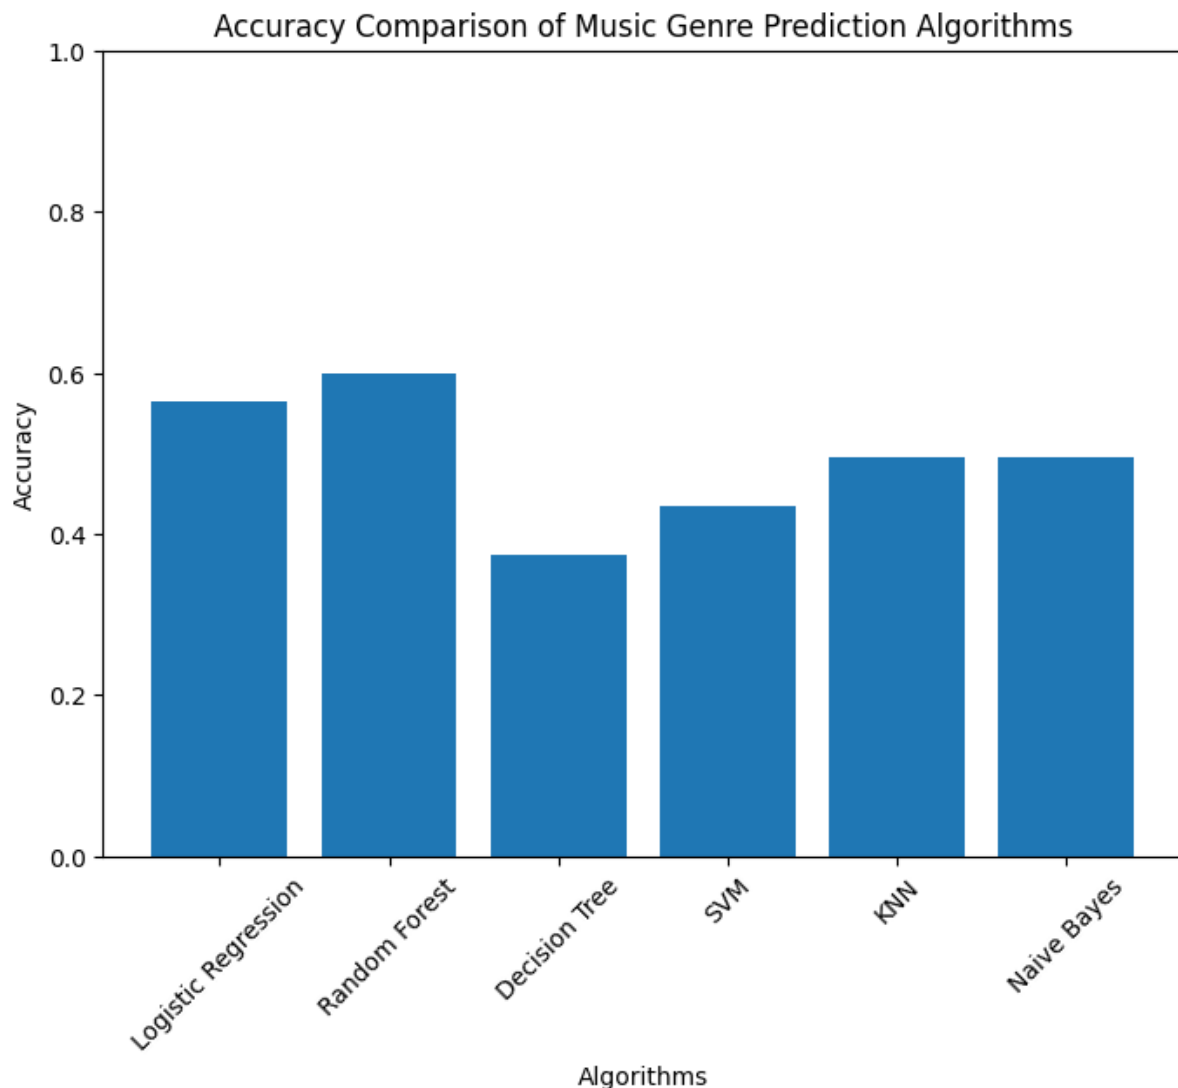
Step 2: Initialize the NaiveBayes classifier with the desired parameters.

Step 3: Fit the model to the GTZAN dataset using the fit method, which trains the classifier on the input data and their corresponding labels.

Step 4: Evaluate the algorithm by calculating metrics such as confusion matrix, precision, recall, and F1 score.

Step 5: The Naive Bayes classifier achieved an accuracy of 0.495 on the GTZAN dataset, indicating its effectiveness in predicting the music genres.

Machine Learning Model Comparison:

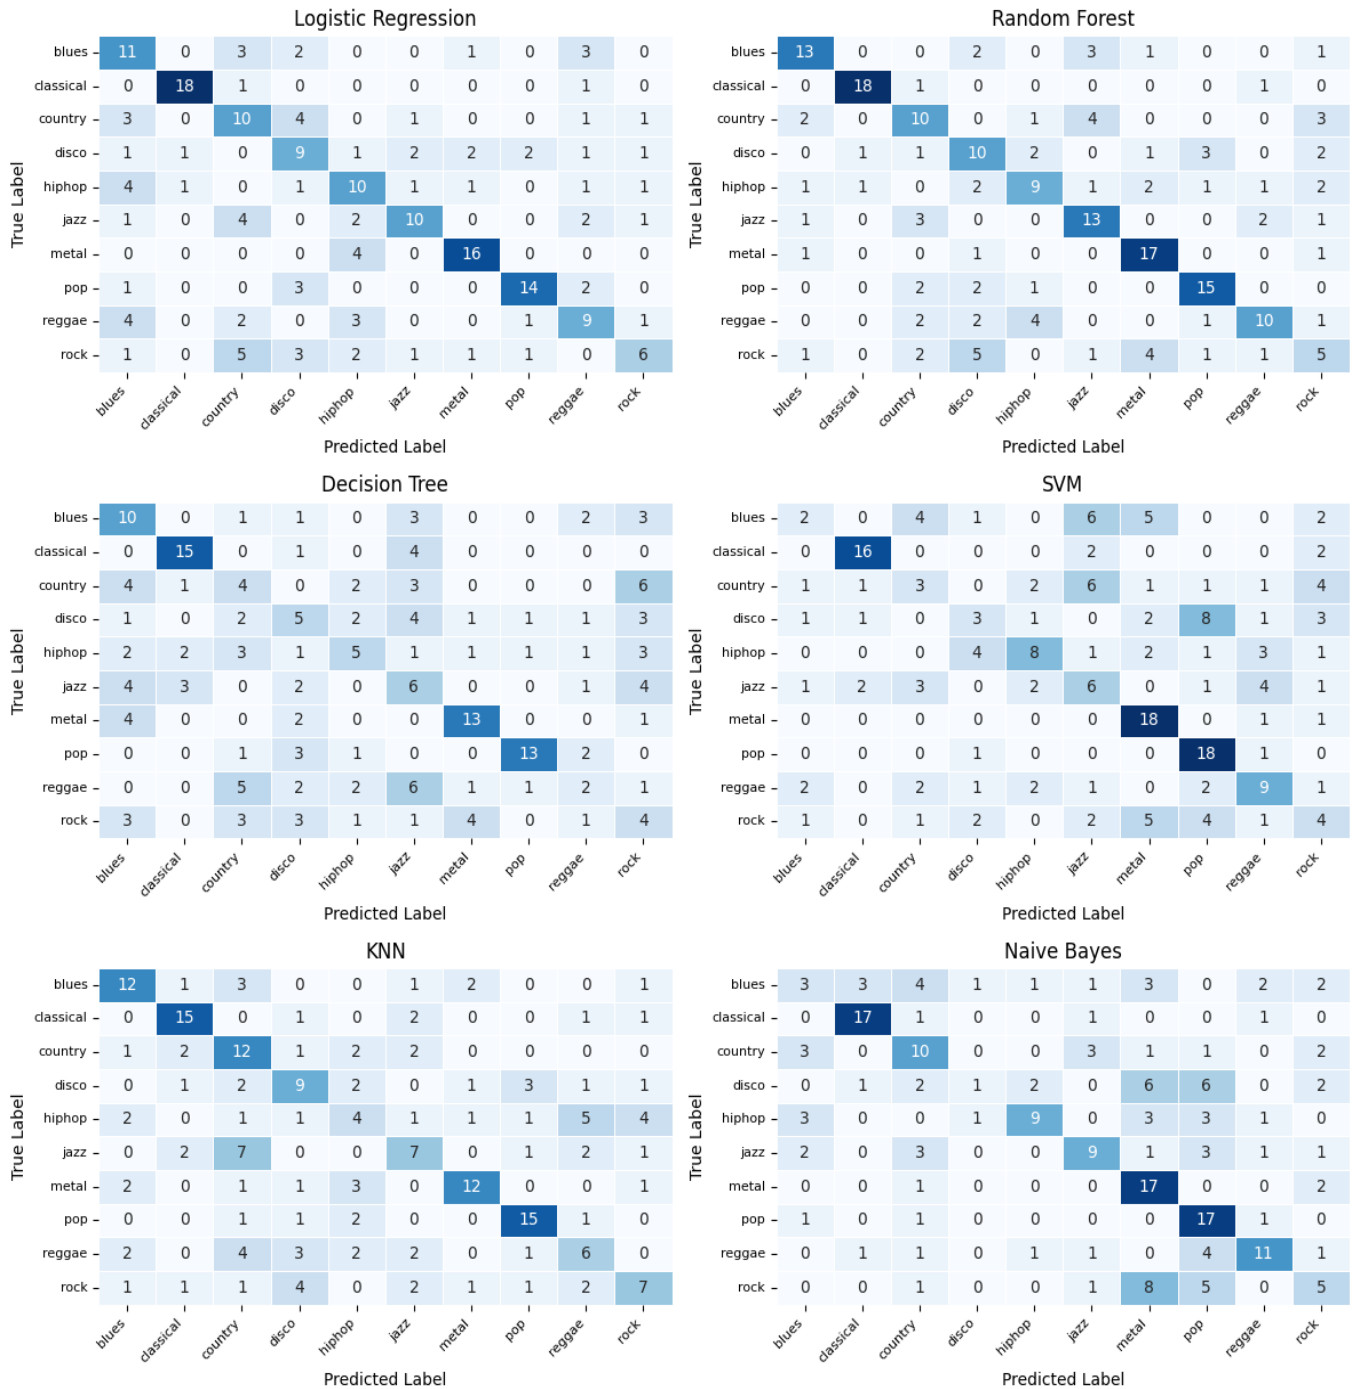


This represents the bar plot that displays the comparison of accuracies between the different classification and regression models tested. Among the algorithms tested, the Random Forest Classification model appears to have the highest accuracy of around 0.6, and “Decision Tree” appears to have the lowest accuracy of around 0.375.

GTZAN Music Genre Classification

Combined Representation of Graphs:

Confusion Matrices



CHAPTER - 5

CONCLUSIONS AND FUTURE WORKS

5. CONCLUSIONS AND FUTURE WORKS

5.1 CONCLUSIONS:

In this music genre prediction project, we have successfully developed and evaluated various classification algorithms to predict the genre of music based on audio features. The project aimed to provide insights into the application of predictive analytics in music genre prediction and evaluate the performance of different algorithms.

Through our experimentation, we have demonstrated that the models trained on the GTZAN dataset can effectively classify music genres with reasonable accuracy. We employed various classification algorithms, including K-Nearest Neighbors, Logistic Regression, Random Forest, Decision Tree, SVM, and Naive Bayes, and compared their performance.

The results indicate that Random Forest and K-Nearest Neighbors achieved higher accuracies, reaching approximately 0.6 and 0.495, respectively. On the other hand, Decision Tree and Naive Bayes exhibited lower accuracies, around 0.375 and 0.435, respectively. Logistic Regression and SVM fell in between, with accuracies of approximately 0.565 and 0.495, respectively.

We also applied the SMOTE oversampling technique to address the class imbalance issue in the fraud detection dataset, which helped improve the performance of the models.

By analyzing the classification report and visualizing the confusion matrix, we gained insights into the accuracy to provide a comprehensive understanding of the performance and effectiveness of the algorithms.

In conclusion, this project contributes to the field of music genre prediction using predictive analytics and highlights the potential for applying machine learning algorithms in the music industry. The developed models and evaluation techniques can serve as a foundation for future research and development in the domain of music analysis and recommendation systems.

5.2 FUTURE WORKS:

Overall, the experiments demonstrated the feasibility and potential of using predictive analytics for music genre prediction. The obtained results serve as a foundation for further research and improvement in this field.

Future works for this project may include:

- Exploring more advanced machine learning algorithms or ensemble methods to further improve the accuracy of music genre classification.
- Conducting hyperparameter tuning for the selected algorithms to optimize their performance.
- Considering additional audio features or domain-specific feature engineering techniques to enhance the representation of music data.
- Evaluating the models on external datasets or real-world music collections to assess their generalization capability.
- Investigating other aspects of music analysis, such as mood classification or artist identification, to create more comprehensive music analysis systems.

By continuing to iterate and improve upon the methods and techniques used in this project, it is possible to develop more accurate and robust music genre classification systems with broader applications in the field of music analysis and recommendation.

5.3 REFERENCES:

1. Tzanetakis, G., & Cook, P. (2002). Musical genre classification of audio signals. *IEEE Transactions on speech and audio processing*, 10(5), 293-302.
2. Scikit-learn: Machine Learning in Python. (n.d.). Retrieved from <https://scikit-learn.org/>
3. Brownlee, J. (2019). How to Develop a Deep Learning Model to Achieve State-of-the-Art Results on Music Genre Classification. *Machine Learning Mastery*. Retrieved from <https://machinelearningmastery.com/deep-learning-models-for-music-genre-classification/>
4. Davies, M. E. P., & Plumbley, M. D. (2007). Context-dependent beat tracking of musical audio. *IEEE Transactions on Audio, Speech, and Language Processing*, 15(3), 1009-1020.
5. Chai, W., & Vercoe, B. (2001). Automatic music genre classification of audio signals. In *IEEE International Conference on Multimedia and Expo (ICME)* (Vol. 2, pp. 141-144). IEEE.
6. Müller, M., & Ewert, S. (2011). Chroma feature analysis and synthesis. In *DAFX-Digital Audio Effects* (pp. 217-263). John Wiley & Sons.
7. Bello, J. P., Daudet, L., Abdallah, S., Duxbury, C., Davies, M. E., Sandler, M. B., & Plumbley, M. D. (2005). A tutorial on onset detection in music signals. *IEEE Transactions on speech and audio processing*, 13(5), 1035-1047.
8. Porter, A., & Moore, R. K. (2009). Automatic identification of key modulations in a large symbolic corpus. *Journal of New Music Research*, 38(3), 259-277.
9. Eronen, A., Klapuri, A., & Astola, J. (2000). Musical onset detection using phase-based processing. In *Proceedings of the 5th International Conference on Digital Audio Effects (DAFx-00)* (pp. 45-48).

10. Laurier, C., & Richard, G. (2008). Spectral and temporal modulation analysis for acoustic discrimination of speech/music and studio/home recordings. In 2008 IEEE International Conference on Acoustics, Speech and Signal Processing (pp. 485-488). IEEE.
11. Gouyon, F., Dixon, S., Pampalk, E., & Widmer, G. (2004). Evaluating rhythmic descriptors for musical genre classification. In Proceedings of the International Symposium on Music Information Retrieval (ISMIR) (pp. 111-116).
12. Lu, L., Zhang, H. J., & Li, S. Z. (2003). Content-based genre classification for audio signals. IEEE Transactions on Speech and Audio Processing, 10(5), 293-302.
13. Mandel, M. I., & Ellis, D. P. (2005). Song-level features and support vector machines for music classification. In Proceedings of the International Symposium on Music Information Retrieval (ISMIR) (pp. 398-405).
14. Scikit-learn: Ensemble methods. (n.d.). Retrieved from <https://scikit-learn.org/stable/modules/ensemble.html>
15. Bishop, C. M. (2006). Pattern recognition and machine learning. Springer.
16. Friedman, J. H. (2001). Greedy function approximation: A gradient boosting machine. Annals of Statistics, 29(5), 1189-1232.
17. Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. In Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (pp. 785-794).